

Department of Computer Engineering

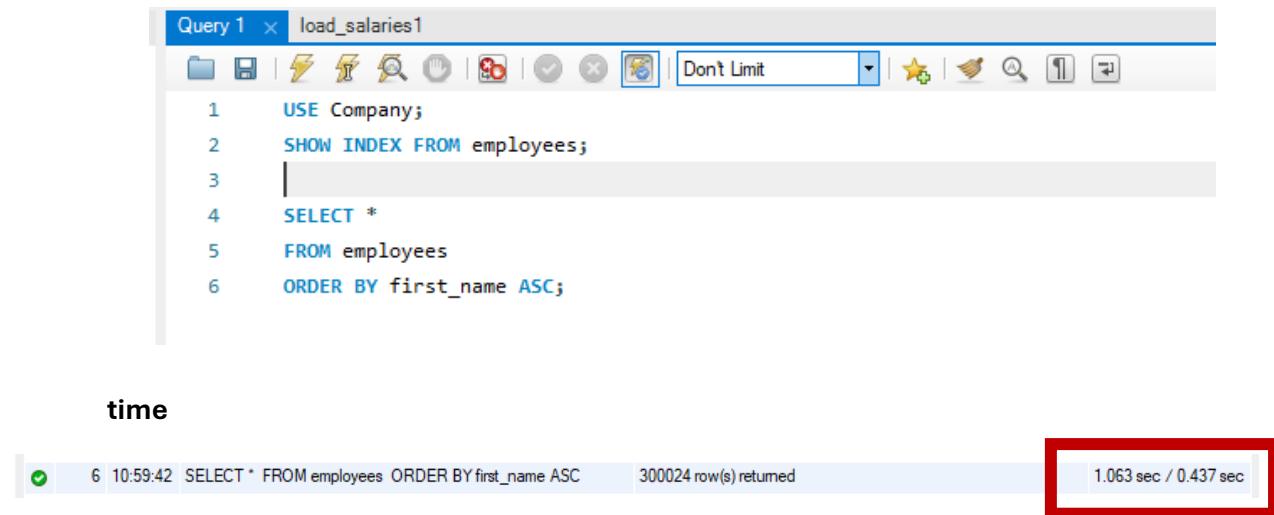
University of Peradeniya

CO527 Advanced Database Systems

Edirimanna Y.H.

E/20/089

1. Assuming no indexes are used, record the query execution time for retrieving all the employees by first name in ascending order.

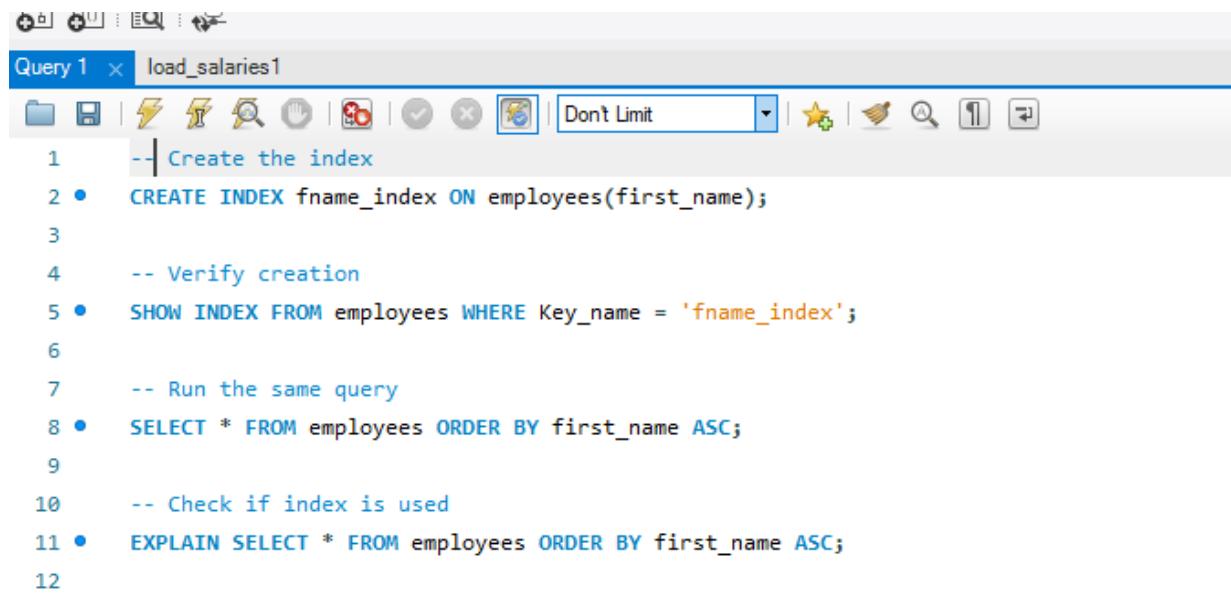


The screenshot shows the MySQL Workbench interface with a query editor window titled "Query 1". The query code is:

```
1 USE Company;
2 SHOW INDEX FROM employees;
3
4 SELECT *
5 FROM employees
6 ORDER BY first_name ASC;
```

Below the query editor, the status bar displays the execution time: "1.063 sec / 0.437 sec".

2. Create an index called fname_index on the first_name of the employee table. Retrieve all the employees by first name and record the query execution time. Observe the performance improvement gained when accessing with index.



```
Query 1 x load_salaries1
+---+---+---+---+---+---+---+---+---+---+---+---+
| 1 | -- Create the index
| 2 | • CREATE INDEX fname_index ON employees(first_name);
| 3 |
| 4 | -- Verify creation
| 5 | • SHOW INDEX FROM employees WHERE Key_name = 'fname_index';
| 6 |
| 7 | -- Run the same query
| 8 | • SELECT * FROM employees ORDER BY first_name ASC;
| 9 |
|10| -- Check if index is used
|11| • EXPLAIN SELECT * FROM employees ORDER BY first_name ASC;
|12|

```

Time

16	11:12:31	SELECT * FROM employees ORDER BY first_name ASC	300024 row(s) returned	0.625 sec / 0.297 sec
----	----------	---	------------------------	-----------------------

Before time is 1.063msec

After indexing time is 0.625msec

3. Which indexing technique has been used when creating the above index?

The screenshot shows the MySQL Workbench interface. The top part is a query editor titled "Query 1" with the session name "load_salaries1". It contains the following SQL code:

```
1 -- Detailed index information
2 • SHOW INDEX FROM employees WHERE Key_name = 'fname_index';
3
4 -- Alternative query
5 • SELECT
6     INDEX_NAME,
7     INDEX_TYPE,
8     COLUMN_NAME,
9     CARDINALITY,
10    SEQ_IN_INDEX
11   FROM information_schema.STATISTICS
12  WHERE TABLE_SCHEMA = 'Company'
13    AND TABLE_NAME = 'employees'
14    AND INDEX_NAME = 'fname_index';
15
```

The bottom part is a "Result Grid" showing the output of the query:

	INDEX_NAME	INDEX_TYPE	COLUMN_NAME	CARDINALITY	SEQ_IN_INDEX
▶	fname_index	BTREE	first_name	1266	1

A sidebar on the right indicates that the current view is "Result Grid".

Type is B Tree indexing

4. Create a unique index on emp_no, first_name and last_name of employees table. Retrieve all the employees by emp_no, first_name and last_name. Observe if there is any performance improvement with respect to question1. If not, explain any possible reason.

```
1
2 • CREATE UNIQUE INDEX emp_composite_index
3   ON employees(emp_no, first_name, last_name);
4
5   -- Test query 1: Using all columns
6 • SELECT * FROM employees
7   WHERE emp_no = 10001
8     AND first_name = 'Georgi'
9     AND last_name = 'Facello';
10
11 • EXPLAIN SELECT * FROM employees
12   WHERE emp_no = 10001
13     AND first_name = 'Georgi'
14     AND last_name = 'Facello';
15
16   -- Test query 2: ORDER BY first_name only
17 • SELECT * FROM employees ORDER BY first_name ASC;
18
19 • EXPLAIN SELECT * FROM employees ORDER BY first_name ASC;
```

There may not be significant performance improvement compared to Question 1 because:

1. The primary key (emp_no) already has a clustered index
2. When ordering by first_name alone, the composite index (emp_no, first_name, last_name) may not be optimal since emp_no is the leading column
3. For the composite index to be effective, queries should use emp_no in WHERE clause or ORDER BY should include emp_no as the first column

5. Take the following 3 queries.

A. select distinct emp_no from dept_manager where from_date>= '1985-01-01' and dept_no>= 'd005';

```
1 •   SHOW INDEX FROM dept_manager;
2
3     -- Query A
4 •   SELECT DISTINCT emp_no
5     FROM dept_manager
6     WHERE from_date >= '1985-01-01'
7     AND dept_no >= 'd005';
8
9
```

emp_no
110511
110567
110725
110765
110800
110854
111035
111133
111400
111534
111692
111784
111877
111939

B. select distinct emp_no from dept_manager where from_date>= '1996-01-03' and dept_no>= 'd005';

Result Grid		 Filter Rows:	<input type="text"/>	Export: 	Wrap Cell Content
	emp_no				
▶	111939				

C. select distinct emp_no from dept_manager where from_date>= '1985-01-01' and dept_no<= 'd009';

The screenshot shows the MySQL Workbench interface. The top window is titled "Query 1" and contains the SQL code for query C. The bottom window is titled "Result Grid" and displays the results of the query, which are employee numbers (emp_no) ranging from 110022 to 111784.

emp_no
110022
110039
110085
110114
110183
110228
110303
110344
110386
110420
110511
110567
110725
110765
110800
110854
111035
111133
111400
111534
111692
111784
.....

I. Choose one single simple index(i.e index on one attribute) that is most likely to speed up all 3 queries giving reasons for your selection.

I chose from_date because:

1. It appears in all three queries with range conditions, from_date has higher cardinality (more distinct values) than dept_no ,Date ranges are more selective than dept_no ranges and MySQL can use this index to filter rows before applying dept_no condition

II. For each of the 3 queries, check if MySQL storage engine used that index. If not, give a short explanation why not. You can prefix your select queries with EXPLAIN EXTENDED or with EXPLAIN to display a query execution plan.

A

```
-- Query A with EXPLAIN
EXPLAIN SELECT DISTINCT emp_no
FROM dept_manager
WHERE from_date >= '1985-01-01'
AND dept_no >= 'd005';
```

Result Grid												
	id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
▶	1	SIMPLE	dept_manager	NULL	index	PRIMARY,dept_no, idx_from_date	PRIMARY	20	NULL	24	58.33	Using where

it shows "PRIMARY": Using the primary key instead.

Reason:

1. from_date = '1985-01-01' is VERY EARLY (near minimum date in table)
 - This condition matches MOST rows in the table (low selectivity)
2. dept_no >= 'd005' eliminates departments d001-d004
 - The PRIMARY KEY includes dept_no, making it useful for this filter
3. Since most rows match from_date condition, MySQL decides:
 - Cost of idx_from_date: scan many index entries + lookup rows
 - Cost of PRIMARY: direct scan with dept_no filtering built-in
 - PRIMARY is more efficient because dept_no is part of the key

B

```
-- Query B with EXPLAIN
EXPLAIN SELECT DISTINCT emp_no
FROM dept_manager
WHERE from_date >= '1996-01-03'
AND dept_no >= 'd005';
```

it shows "idx_from_date": The index is being used

1. from_date = '1996-01-03' is MUCH LATER (closer to maximum date)
 - This condition is HIGHLY SELECTIVE (matches fewer rows)
 - Most managers started before 1996, so this filters out many rows
2. The idx_from_date can quickly skip to 1996 entries

3. Since from_date filters out most rows efficiently:

- Cost of idx_from_date: scan few index entries + lookup few rows
- Cost of PRIMARY: must scan more entries to filter
- idx_from_date is more efficient due to high selectivity

C

```

Query 1 | salary_audit | company_titles
| File | Edit | View | Tools | Help | Don't Limit | 
1 -- Query C with EXPLAIN
2 EXPLAIN SELECT DISTINCT emp_no
3 FROM dept_manager
4 WHERE from_date >= '1985-01-01'
5 AND dept_no <= 'd009';

```

	id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
▶	1	SIMPLE	dept_manager	NULL	index	PRIMARY,dept_no,`idx_from_date`	PRIMARY	20	NULL	24	100.00	Using where

it shows "PRIMARY": Using the primary key instead.

1. from_date = '1985-01-01' matches MOST rows (low selectivity)
2. dept_no <= 'd009' matches ALL or MOST departments
 - If there are only 9 departments (d001-d009), this matches everything
 - Very low selectivity for dept_no condition
3. PRIMARY KEY scan is efficient because:
 - Both conditions have low selectivity
 - No advantage to using idx_from_date
 - PRIMARY provides better data locality (clustered index)

6.n Consider the queries you wrote for questions 2 - 10 in PreLab assignment. Give with short explanations, which attributes on which relations should be used for creating indexes that could speed up your queries.

Query #	Query Description	Tables Involved	Recommended Index(es)	Column(s)	Reasoning
2	Top 10 family names	employees	None	-	Requires full table scan for GROUP BY on all last_names. Index won't help when aggregating all distinct values.
3	Number of Engineers per department	departments, dept_emp, titles	idx_title idx_dept_emp_todate	title to_date	• Filter titles containing 'Engineer'

Query #	Query Description	Tables Involved	Recommended Index(es)	Column(s)	Reasoning
4	Female managers who were senior engineers	employees, dept_manager, titles	idx_gender idx_title	gender title	<ul style="list-style-type: none"> Filter current employees (to_date = '9999-01-01') Equality filter on gender = 'F' Equality filter on title = 'Senior Engineer'
5	Employees with salary > 115000	employees, dept_emp, departments, titles, salaries	idx_salary idx_salaries_todate idx_dept_emp_todate	salary to_date to_date	<ul style="list-style-type: none"> Range query on salary > 115000 Filter current salaries Filter current employees
6	Senior employees (>50 years, >10 years service)	employees	idx_birth_date idx_hire_date	birth_date hire_date	<ul style="list-style-type: none"> Help with date range calculations Though TIMESTAMPDIFF function prevents direct index usage, indexes still assist with range scans
7	Employees NOT in HR	employees, dept_emp, departments	idx_dept_name idx_dept_emp_todate	dept_name to_date	<ul style="list-style-type: none"> Filter by department name Filter current employees
8	Employees earning more than all Finance employees	employees, salaries, dept_emp, departments	idx_salary idx_dept_emp_composite	salary (dept_no, to_date)	<ul style="list-style-type: none"> Salary comparison with MAX() Composite index for department filtering
9	Employees earning more than company average	employees, salaries, dept_emp, departments	idx_salary_todate	(salary, to_date)	<ul style="list-style-type: none"> Composite index for salary comparison Filter current salaries efficiently
10	Salary difference: Senior Engineers vs All	salaries, titles	idx_title idx_salaries_todate	title to_date	<ul style="list-style-type: none"> Filter by title = 'Senior Engineer' Filter current salaries for AVG calculation

7. Assume that most of the queries on a relation are insert/update/delete. What will happen to the query execution time if that relation has an index created?

If a relation has indexes and most queries are INSERT/UPDATE/DELETE operations, the query execution time will **INCREASE**

- With indexes: Fast SELECT, Slow INSERT/UPDATE/DELETE
- Without indexes: Slow SELECT, Fast INSERT/UPDATE/DELETE
- Ex –
- Table with 5 indexes: - 1 INSERT operation = 1 table write + 5 index updates = 6 total operations - Without indexes = 1 table write = 1 operation Result: 6x slower for writes

For write-heavy (OLTP) systems, minimize indexes and only create essential ones. For read-heavy (OLAP/reporting) systems, more indexes are acceptable as the read performance gain outweighs the write penalty.