

Department of Computer Engineering

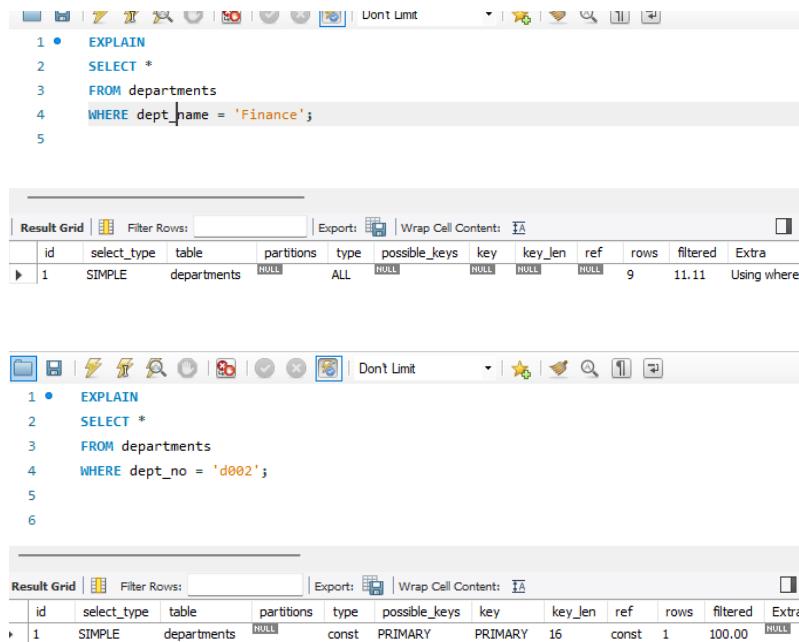
University of Peradeniya

CO527 Advanced Database Systems

Edirimanna Y.H.

E/20/089

1) Troubleshooting Simple Queries Using EXPLAIN



```
1 • EXPLAIN
2 SELECT *
3 FROM departments
4 WHERE dept_name = 'Finance';
5
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	departments	NULL	ALL	NULL	NULL	NULL	NULL	9	11.11	Using where


```
1 • EXPLAIN
2 SELECT *
3 FROM departments
4 WHERE dept_no = 'd002';
5
6
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	departments	NULL	const	PRIMARY	PRIMARY	16	const	1	100.00	NULL

Query 1: dept_name = 'Finance'

- key = NULL → No index is used
- type = ALL → Full table scan
- rows → MySQL checks all rows in departments
- Reason: dept_name is not indexed

Interpretation: MySQL scans the entire department's table to find matching records, which is inefficient and does not scale well for large datasets.

Query 2: dept_no = 'd002'

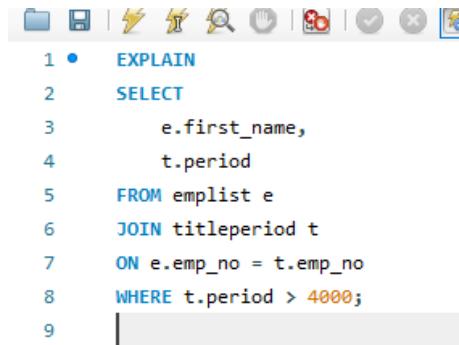
- key = PRIMARY → Primary key index is used
- type = const → Fastest possible access type
- rows = 1 → Only one row examined
- Reason: dept_no is the primary key

Interpretation: MySQL directly locates the required row using the primary key index, resulting in minimal row access and optimal performance.

2) Join Query Analysis Using EXPLAIN

```
1 • CREATE TABLE emplist
2   SELECT emp_no, first_name FROM employees;
3
4 • CREATE TABLE titleperiod
5   SELECT emp_no, title, DATEDIFF(to_date, from_date) AS period
6   FROM titles;
7
```

Initial Query



The screenshot shows a MySQL query editor window. At the top, there are several icons: a folder, a file, a lightning bolt, a magnifying glass, a search icon, a hand, a red circle with a question mark, a checkmark, a close button, and a refresh button. Below the toolbar, the text area contains the following SQL code:

```
1 • EXPLAIN
2   SELECT
3     e.first_name,
4     t.period
5   FROM emplist e
6   JOIN titleperiod t
7   ON e.emp_no = t.emp_no
8   WHERE t.period > 4000;
9
```

EXPLAIN Output Analysis

- Join type: ALL (full table scan on both tables)
- Rows examined: emplist_rows × titleperiod_rows
- MySQL uses a nested loop join
- Large number of row combinations must be checked
- No indexes available to speed up matching

Estimated Row Combinations

- emplist has 300,024 rows
- titleperiod has 443308 rows

Then MySQL may evaluate hundreds of millions of row combinations, leading to poor performance. This is why the query is extremely slow.

3) Indexing for Optimization

```
1 •  ALTER TABLE emplist
2      ADD PRIMARY KEY (emp_no);
3
4 •  CREATE INDEX empno_idx
5      ON titleperiod(emp_no);
6
```

Before Indexing

```
25 11:05:13 CREATE TABLE emplist SELECT emp_no,first_name FROM employees          300024 row(s) affected Records: 300024 Duplicates: 0 Warnings: 0
26 11:05:15 CREATE TABLE titleperiod SELECT emp_no,title,DATEDIFF(to_date,from_date) AS period FROM titles        443308 row(s) affected Records: 443308 Duplicates: 0 Warnings: 0
                                                2.50 sec
                                                4.032 sec
```

After Indexing

```
30 11:54:22 EXPLAIN SELECT e.first_name, t.period FROM emplist e JOIN titleperiod t ON e.emp_no = t.emp_no WHERE t.period > 4000      2 row(s) returned
                                                0.015 sec / 0.000 sec
```

The Improvement

Before Indexing

Full table scans

Huge row combinations

key = NULL

Slow

After Indexing

Index lookups

Drastically reduced

key = PRIMARY / idx_emp_no

Fast

Is it possible to optimize the query execution further?

Yes it is possible to optimize. We can use Composite Index for this. This allows efficient join and efficient filtering.

```
CREATE INDEX idx_emp_period
ON titleperiod(emp_no, period);
```

4) Query Rewriting Techniques

FORCE INDEX

```

1 • SELECT first_name
2   FROM empelist FORCE INDEX(PRIMARY)
3   WHERE emp_no > 1000;
4
5

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

first_name
Georgi
Bezalel
Parto
Christian
Kyochi
Anneke
Tsvetan
Saniya
Sumant
Duangkaew
Mary
Patrício
Eberhardt
Berni
Guoxiang
Kazuhito
Cristinel

IGNORE INDEX

```

1 • EXPLAIN
2   SELECT *
3   FROM titleperiod IGNORE INDEX (empno_idx)
4   WHERE period > 4000;
5

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: |

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	titleperiod	NULL	ALL	NULL	NULL	NULL	NULL	442407	33.33	Using where

STRAIGHT_JOIN

```

1 • SELECT STRAIGHT_JOIN e.first_name, t.period
2   FROM empelist e
3   JOIN titleperiod t
4   ON e.emp_no = t.emp_no
5   WHERE t.period > 4000;
6

```

Result Grid | Filter Rows: | Export: | Wrap Cell Content: | Fetch rows: |

first_name	period
Georgi	2926512
Bezalel	2922821
Parto	2923065
Christian	2923067
Kyochi	2922781
Anneke	2925011
Tsvetan	2922995
Sumant	2923353
Duangkaew	2922708
Patrício	2921223
Eberhardt	2926761
Berni	2923769
Kazuhito	2922264
Cristinel	2921360
Kazuhide	2923309
Lilian	2921821
.....

These techniques provide **manual control** when the optimizer's choice is suboptimal.